

CS2107

Introduction to Information Security

AY2022/23 Semester 2

Notes by Jonathan Tay

Last updated on January 26, 2023

Contents

I	Cryptography	1
1	Encryption	1
1.1	Classical Ciphers	1
1.2	Modern Ciphers	1
1.3	Attacks	2
II	Authentication	3
2	Passwords	3
2.1	Attacks	3
2.2	n -Factor Authentication	3
3	Biometrics	4
3.1	Attacks	4

Part I

Cryptography

1 Encryption

Symmetric-key encryption uses the same key k for encryption and decryption.

Any plaintext x that is encrypted to produce a ciphertext $c = E_k(x)$ must also be decryptable to recover the plaintext $D_k(E_k(x)) = x$ — (**correctness**).

It must also be difficult or impossible to derive useful information of the key or the plaintext from the ciphertext — (**security**).

attack models: information

Given access to an **oracle** with varying levels of information, adversaries can conduct several models of attack:

- **ciphertext only (COA)**: only ciphertexts and properties of the plaintext are known
- **known-plaintext (KPA)**: ciphertexts of corresponding plaintexts are known
- **chosen-plaintext (CPA)**: ciphertexts can be derived from arbitrary plaintexts
- **chosen-ciphertext (CCA2)**: plaintexts can be derived from arbitrary ciphertexts

attack models: goals

Adversaries may have varying goals:

- **total break**: find the encryption key
- **partial break**: decrypt a ciphertext or obtain information about a plaintext
- **indistinguishability**: distinguish ciphertexts of different plaintexts better than random chance

1.1 Classical Ciphers

substitution cipher

Construct a **substitution table** used as the key from some permutation of all possible characters used in the plaintext.

Replace each character in the plaintext with the corresponding character in the substitution table to produce the ciphertext.

Substitution ciphers are vulnerable to various attack models:

- **general**: brute-force exhaustive search on all keys
- **KPA**: substitution table reconstruction by comparison of characters
- **COA**: frequency analysis of characters between ciphertext and plaintext

permutation ciphers

Group characters of the plaintext into blocks of equal size, then rearrange the characters in each block according to a permutation key.

Extremely vulnerable under KPA and COA attacks as the permutation key can be easily reconstructed, especially if the plaintext language is known.

Note that applying substitution or permutation ciphers several times alone in succession is equivalent to applying it once — the ciphertext is not any more secure!

one-time pad

Generate a key of equal bit length as the plaintext, then XOR them together to encrypt. XOR the ciphertext with the key to decrypt:

correctness:

1. $c = x \text{ XOR } k$
2. $(x \text{ XOR } k) \text{ XOR } k = x \text{ XOR } (k \text{ XOR } k) = x \text{ XOR } 0 = x$

security: Even if an adversary obtains a ciphertext-plaintext pair and derives the key (trivial), the key is useless as it is not re-used.

Exhaustive search does not work on OTPs because no meaningful information about the plaintext can be learnt (**perfect secrecy**) — every decrypted plaintext is equally probable.

1.2 Modern Ciphers

DES (data encryption standard) has a key length of 56 bits, and as such is vulnerable to exhaustive search today.

AES (advanced encryption standard) has a block length of 128 bits, and a key length of 128, 192, or 256 bits.

Both DES and AES are **block ciphers** which inputs and outputs of fixed length. This means plaintext must be padded and divided into blocks before encryption.

There are several methods of extending an encryption from single to multiple blocks.

ECB (electronic code book) mode

Encrypt each block independently, then concatenate the ciphertexts.

If the encryption scheme is *deterministic*, then any two blocks of plaintext encrypts to the same ciphertext, which leaks information.

AES is inherently deterministic, and as such requires an **initialization vector (IV)** for probabilistic encryption.

An IV is an arbitrary plaintext value (i.e. not secret) that should be unique for each encryption:

1. Let X and Y be two different plaintexts, U and V their ciphertexts, and K the IV.
2. $U \oplus V = (X \oplus K) \oplus (Y \oplus K)$.
3. By associativity and commutativity, $U \oplus V = (X \oplus Y) \oplus (K \oplus K) = X \oplus Y$.
4. $X \oplus Y$ can leak information.
5. Hence, we need a unique IV for each encryption.

The IV is also needed for decryption and as such is sent together with the ciphertext.

CBC (cipher block chaining) mode

XOR the first block with the IV, then encrypt it.

Subsequent blocks are XOR'd with the previous ciphertext block before encryption.

CTR (counter) mode

Generate a cryptographically-secure pseudo-random sequence from the secret key and IV *and then* encrypt it (not the plaintext).

Plaintext blocks can be streamed in are XOR'd with the encrypted sequence (**stream cipher**).

Decryption uses the same process.

1.3 Attacks**meet-in-the-middle attack (MITM) (\in KPA)**

Suppose we apply DES twice sequentially using keys K_1 and K_2 to encrypt a plaintext p into ciphertext $c = E_{K_2}(E_{K_1}(p))$.

Recall that in a KPA, an adversary aiming for a total break wants to find K_1 and K_2 given c and p .

Define the intermediate encryption value $x = E_{K_1}(p)$. For decryption, we also have $x' = D_{K_2}(c)$.

Now, the adversary can pre-compute and store all 2^{56} possible values of x and all 2^{56} possible values of x' in two hash tables.

With 2^{57} possible pairs of (x, x') , exhaustive search will yield the keys K_1 and K_2 .

Hence, instead of the expected key strength of $2 \times 56 = 112$, by trading space for time, the key strength is only 57.

Any sequential encryption scheme is vulnerable to MITM.

Since MITM only reduces key strength, it can be remedied by increasing the number of rounds in the encryption scheme, e.g. **triple DES** (3DES).

padding oracle attack (\in CCA2)

Suppose AES CBC with **PKCS#7 padding** is used to encrypt a plaintext p into ciphertext c with 3 blocks of size of 4 bytes, giving a total of 12 bytes.

PKCS#7 padding

If a plaintext does not have a length that is exactly a multiple of the block size, let l be the number of bytes needed to fill (pad) the last block.

The last l bytes of the plaintext are set to l .

Otherwise, an additional block is added with all bytes set to 0.

An adversary wants a partial break, i.e. decrypt c to get p , with access to a **padding oracle** which returns true if the padding is valid, and false otherwise.

To decrypt AES CBC, each block of c is first decrypted with the key, then XOR'd with the previous block (or the IV for the first block).

This means that a change in any byte of c , say c_8 , will result in a different value for the plaintext byte in the same position of that and every subsequent block, i.e. p_8 and p_{12} .

Therefore, by exhausting all 256 possible values of c_8 , the padding oracle will return true for the one correct value of $p'_{12} = D_k(c_{12}) \oplus c'_8$.

Since $p'_{12} = p_{12} \oplus c_8 \oplus c'_8$, and the values of p'_{12} , c_8 and c'_8 are known, $p_{12} = p'_{12} \oplus c_8 \oplus c'_8$!

forcing a specific value x for p'_i

To proceed further with the attack on p_{11} via c_7 , we need to force the value of p'_{12} to be $0x2$, so that we can brute-force c'_7 to get p'_{11} to be $0x2$.

Hence, we want a value for c''_8 which achieves this, which is $c_8 \oplus p_{12} \oplus x$, where $x = 0x02$.

We can repeat this process for p_{11} after setting $c_8 = c''_8$, and so on, truncating c so that the padding does not exceed one block.

Any encryption scheme using CBC is vulnerable to the padding oracle attack, and so is CTR if it uses padding.

Part II

Authentication

Authentication is the process of assuring that a communicating entity or the origin of a piece of information is that which it claims to be.

Authenticity implies integrity.

2 Passwords

Passwords are secrets shared between two entities, but unlike secret keys, passwords are human-generated and memorizable.

bootstrapping

The establishment of a user's password with a server, stored in a file together with their identity.

A user can choose a password and send it to the server, or, the server can choose a default password. A secure communication channel is required.

authentication

Password verification can take place *interactively* where the user is prompted for their password, or *non-interactively* where the password is sent to the server in plaintext.

Access is granted if and only if the password is correct.

However, weak authentication methods are vulnerable to **replay attacks** where an attacker can sniff information from the channel and replay it to impersonate the user.

password reset

Systems typically link an account to a recovery email address, allowing the user to reset their password via a link containing a one-time-password (OTP).

Ownership of the email address proves the entity is authentic, as this is 2FA.

2.1 Attacks

Default password stealing or password interception (e.g. mail) can occur during bootstrapping.

Passwords can also be stolen by **shoulder-surfing**, **sniffing** over unencrypted networks, **keylogging**,

phishing (social engineering, including spear-phishing), **cache-exploitation**, or **insider attacks**.

Regular training, phishing exercises, and blacklisting of known phishing sites can combat phishing attacks.

dictionary attacks

Dictionary attacks are a brute-force attack where an attacker uses a pre-selected collection of phrases to guess the password.

During an **online dictionary attack**, the attacker must interact with the authentication server, which acts as an oracle.

Severs can combat this by limiting the number of attempts or introducing delays between attempts.

In an **offline dictionary attack**, the attacker has to first obtain some information about the password (e.g. an encrypted database) before searching through it without interacting with the server (e.g. a password-protected PDF).

side-channel attacks

Side channel attacks rely on information leaked by the system, into the environment which can be intercepted by an attacker and used to derive the password.

For example, sound and timing of keyboard presses.

2.2 n-Factor Authentication

Examples of **factors**:

1. **something you know**: password, PIN
2. **something you have**: hardware token, smartphone, ATM card
3. **who you are**: biometrics

one-time-password (OTP) tokens

Some form of hardware installed with a secret key shared with the server that is used to generate an OTP for a specific and limited event (e.g. login, transaction).

The password generation process can be either **time-based** or **sequence-based** (e.g. a counter for button presses).

Because OTPs are limited in their validity, they cannot be reused and are therefore not vulnerable to replay attacks.

When comparing between possible hardware, e.g. a physical token or a smartphone app, they will each have their own security trade-offs, i.e. there is no single best solution.

3 Biometrics

Biometrics are physical characteristics of a user, e.g. fingerprints.

Unlike passwords, they can't be changed, shared, or forgotten, but they introduce a probability of **false matches** and **false non-matches**.

The **false match rate** (FMR) is the probability that a false biometric capture is incorrectly accepted as belonging to a user.

The **false non-match rate** (FNMR) is the probability that a genuine biometric capture is incorrectly rejected as not belonging to a user.

FMR and FNMR are inversely related — a lower threshold for accepting a biometric capture increases the FMR and decreases the FNMR, and vice versa.

3.1 Attacks

Biometric systems are vulnerable to **spoofing** where an attacker uses a fake biometric capture to impersonate a user.

To combat this, biometric systems can use **liveness detection** where the system checks that the user is present and not an image, recording, etc..